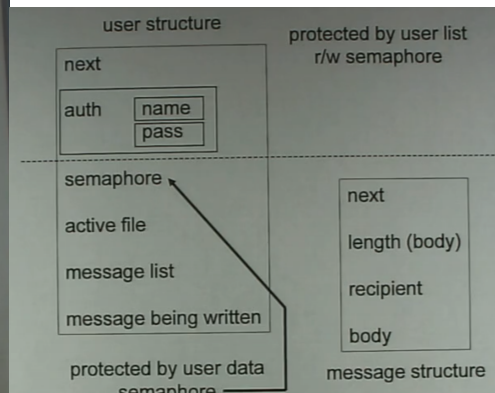
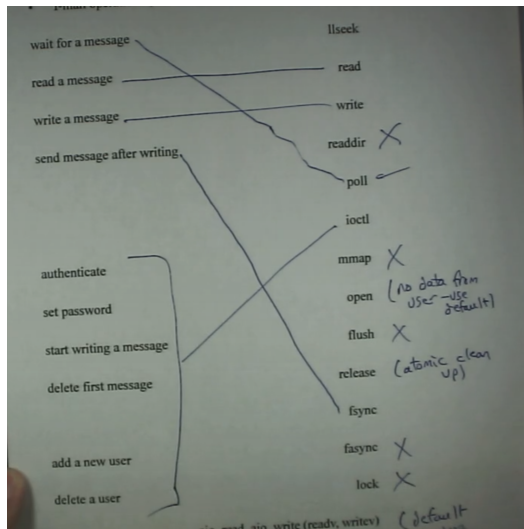


- Block devices (underlies file systems)

- data accessible only in blocks of fixed size, with size determined by device
- can be addressed randomly
- transfers to/from device are usually buffered (to) and cached (from) for performance
- examples: disks, CD ROM, DVD

- Character device:

- contiguous space (or spaces) of bytes
- some allow random access (e.g., magnetic tape driver)
- others available only sequentially (e.g., sound card)
- examples: keyboard, terminal, printers



- What can block?

- read — wait for a new message
- poll — wait for a new message
- write — typically needs to wait for the device, but not in I-mail
- (so only readers/pollers block in I-mail)


- When are readers (or pollers) awoken?

- new message delivered
- user is deleted (nothing is coming now!)

- Protocol for sleeping

- release all locks
- check conditions without locks
- go to sleep
- when awoken
 - reacquire locks
 - recheck all validity requirements for waking (otherwise go back to sleep; false alarm)

- a few small items → **kmalloc**
- a lot of items, repeatedly → **slab cache**
- a big, physically contiguous region → **free pages**
- a big area of virtual memory → **vmalloc** (not necessarily physically contiguous)



- Similarities

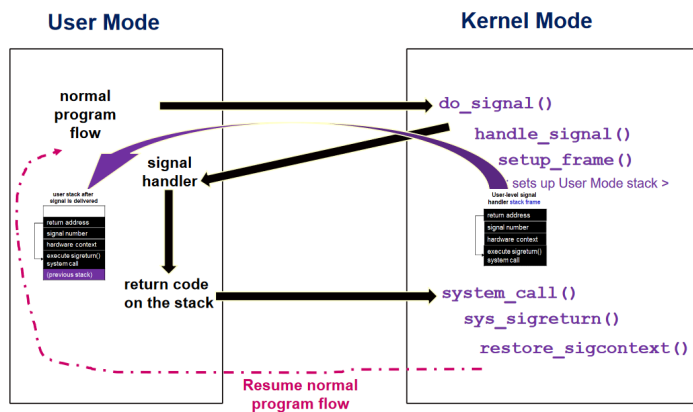
- asynchronous w.r.t. program execution
- not queued — like a physical line; sending twice may cause handler to execute once or twice
- can be ignored, blocked, or caught
- but has some NMIs (SIGKILL, SIGSTOP)
- handler can be changed from default
- only data given to handler is signal # (traditional model)
- signal is (by default) blocked while handler executes

- Differences

- generated by software (kernel or a program via a system call)
- no "device" associated with signal; only software with permission can send a signal (single permission allows any signal to be sent)
- further differences with POSIX (Portable Operating System Interface) model (as opposed to traditional model)

When are signals delivered? (see entry.S)

- check **sigpending** (in task structure) when returning from
 - any interrupt
 - any exception
 - any system call
- only deliver to currently executing process



Permission check for signal generation

- sysadmin (or kernel, or privileged call) can always send
- process with same user id can always send
- process with same login session can send SIGCONT

If generated signal is not being ignored

- it is added to pending signals
- a sleeping recipient is then woken up (kicked out of wait queues, for example)

